

COURSE "AUTOMATED PLANNING: THEORY AND PRACTICE"

CHAPTER 12: THE RELAXATION PRINCIPLE: A CLOSER LOOK

Teacher: **Marco Roveri** - `marco.roveri@unitn.it`
M.S. Course: Artificial Intelligence Systems (LM)
A.A.: 2023-2024
Where: DISI, University of Trento
URL: `https://bit.ly/3zOkGk8`



Last updated: Sunday 5th November, 2023

TERMS OF USE AND COPYRIGHT

USE

This material (including video recording) is intended solely for students of the University of Trento registered to the relevant course for the Academic Year 2023-2024.

SELF-STORAGE

Self-storage is permitted only for the students involved in the relevant courses of the University of Trento and only as long as they are registered students. Upon the completion of the studies or their abandonment, the material has to be deleted from all storage systems of the student.

COPYRIGHT

The copyright of all the material is held by the authors. Copying, editing, translation, storage, processing or forwarding of content in databases or other electronic media and systems without written consent of the copyright holders is forbidden. The selling of (parts) of this material is forbidden. Presentation of the material to students not involved in the course is forbidden. The unauthorised reproduction or distribution of individual content or the entire material is not permitted and is punishable by law.

The material (text, figures) in these slides is authored by Jonas Kvarnström and Marco Roveri.

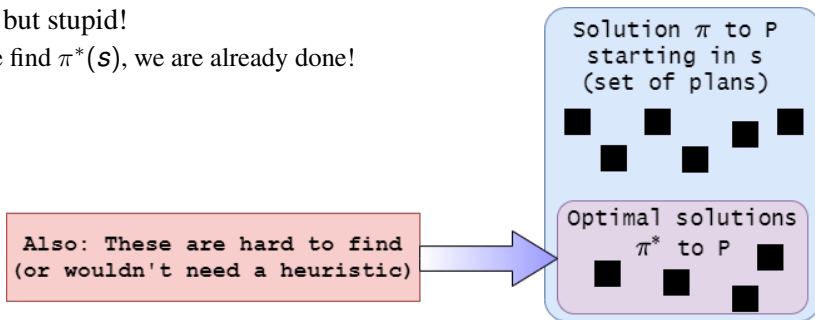
THE PROBLEM

- We have
 - An arbitrary planning problem $P = \langle \Sigma, s_0, S_g \rangle$
- Suppose we want:
 - A way to compute **admissible heuristics** $h(s)$
 - Given P and some states s in the search space

What do we do?
Where do we start?
How do we think?

FUNDAMENTAL IDEAS

- One obvious method: Every time we need $h(s)$ for some state s ...
 - 1 Solve P optimally starting from s , resulting in an *actual* solution $\pi^*(s)$
 - 2 Let $h(s) = h^*(s) = \text{cost}(\pi^*(s))$
 - Admissible - why?
- Obvious, but stupid!
 - If we find $\pi^*(s)$, we are already done!



FUNDAMENTAL IDEAS (CONT.)

- Let's modify the obvious idea:

- Change/Transform P to make it *easy* (quick) to solve

- But make sure optimal solution cannot become more expensive!

- Example: Add more goal states to S_g
 \implies more ways to reach them!

Relaxation will be **one specific way** of (1) **finding** a simplified transformation, and (2) **proving** "not more expensive"!

- Compute an admissible heuristic:

- Solve the modified planning problem optimally

- $h(s)$ = cost of optimal solution for modified problem

$$\leq$$

- $h^*(s)$ = cost of optimal solution for original problem

- Definition of admissibility!


- Preferably

- Keep $h(s)$ as close as possible to $h^*(s)$ - we want *strong cost information*!

FUNDAMENTAL IDEAS (CONT.)

- More formally:

- Before planning, **find** a **simpler** problem P' , such that in every state s (of P):
 - We can **quickly** transform s into a state s' for P'
 - We can **quickly** find an optimal solution π' for P' starting in s'
 - The solution is **never more expensive**: $cost(\pi') \leq cost(\pi^*)$





FUNDAMENTAL IDEAS (CONT.)

- During planning:
 - Every time we need $h(s)$ for some state s :
 - Transform s into s'
 - Quickly solve problem P' optimally starting in s' , resulting in solution π' - for the *transformed* problem
 - Let $h(s) = cost(\pi')$
 - Throw away π' : It isn't interesting itself
- We then know:
 - $h(s) = cost(\pi'(s')) = cost(\text{optimal-solution}(P')) \leq cost(\text{optimal-solution}(P))$
 - $h(s)$ is admissible!

FUNDAMENTAL IDEAS (CONT.)

- Important:

- What we **need**: $cost(\text{optimal-solution}(P')) \leq cost(\text{optimal-solution}(P))$
- Could** use **any** transformation, even with **completely disjoint** solution sets, **if** we just have a **proof** that optimal solution to P' are not more expensive!




Difficult to find transformations, prove correctness - we need a *method*!

FUNDAMENTAL IDEAS (CONT.)


- How to prove $\text{cost}(\text{optimal-solution}(P')) \leq \text{cost}(\text{optimal-solution}(P))$?
 - Sufficient criterion:** One optimal solution to P remains a solution for P'
 - $\text{cost}(\text{optimal-solution}(P')) = \min\{\text{cost}(\pi) \mid \pi \text{ is any solution to } P'\} \leq \text{cost}(\text{optimal-solution}(P))$

Includes the optimal solution to P ,
so $\min\{\dots\}$ cannot be greater



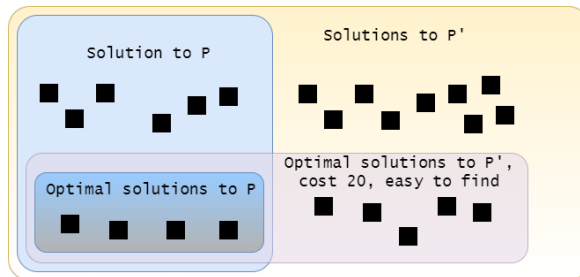
FUNDAMENTAL IDEAS (CONT.)

- Another sufficient criterion: **All solutions** to P **remain** solutions for P'
 - Stronger, but often **easier to prove**
 - **This** is called **relaxation**: P' is a relaxed version of P
 - **Relaxes** the constraints on what is accepted as a solution:
The **is-solution(plan)** test is "expanded, relaxed" to cover additional plans



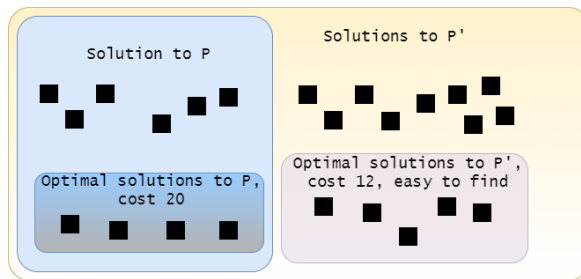
FUNDAMENTAL IDEAS (CONT.)

- Case I: P' has identical cost (for some starting state S)
 - Unlikely!




FUNDAMENTAL IDEAS (CONT.)

- Case II: P' has lower cost (for some starting state S)

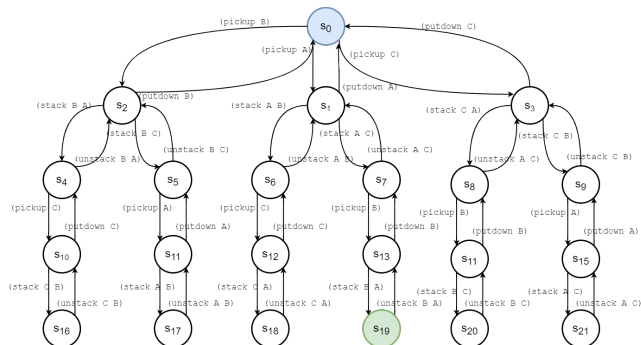


RELAXATION FOR PLANNING PROBLEMS

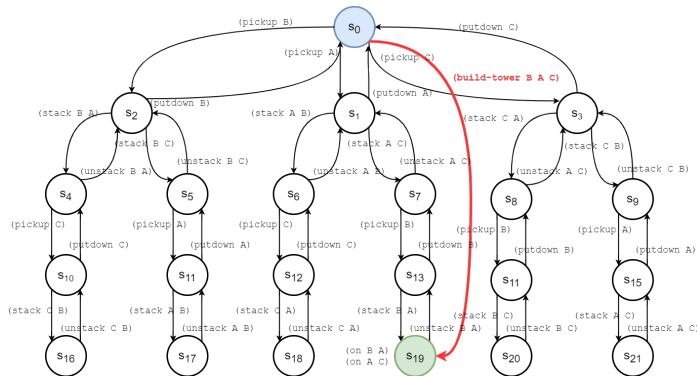
- A classical planning problem $P = \langle \Sigma, s_0, S_g \rangle$ has a **set of solutions**
 - $Solutions(P) = \{\pi | \pi \text{ is an executable action sequence leading from } s_0 \text{ to some state in } S_g\}$
- Suppose that:
 - $P = \langle \Sigma, s_0, S_g \rangle$ is a classical planning problem
 - $P' = \langle \Sigma', s'_0, S'_g \rangle$ is another classical planning problem
 - $Solutions(P) \subseteq Solutions(P')$
- Then (and only then): P' is a relaxation of P !



- A simple planning problem (domain + instance)
 - Blocks world 3 blocks
 - Initially all blocks on the table
 - Goal: $(\text{and } (\text{on } B \ A) \ (\text{on } A \ C))$ (only satisfied in s_{19})
 - Solutions: **All** paths from init to goal (infinitely many - can have cycles)




- Adding new actions
 - All old solutions still valid, but new solutions may exist
 - Modified the STS by adding new edges/transitions
 - This particular example: shorter solutions appear!



RELAXATION EXAMPLE (CONT.)

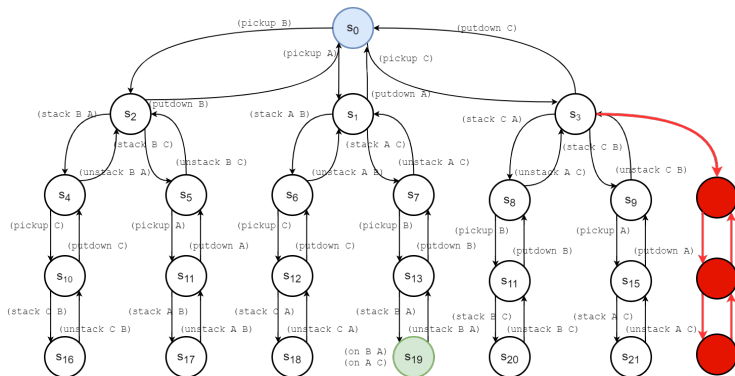
Adding new actions

- In other cases, the new actions may not "help"!
- New solutions ($s_0 \rightarrow s_{19} \rightarrow s_{20}$) are *longer* as well as *more expensive*
- Still relaxation!



RELAXATION EXAMPLE (CONT.)

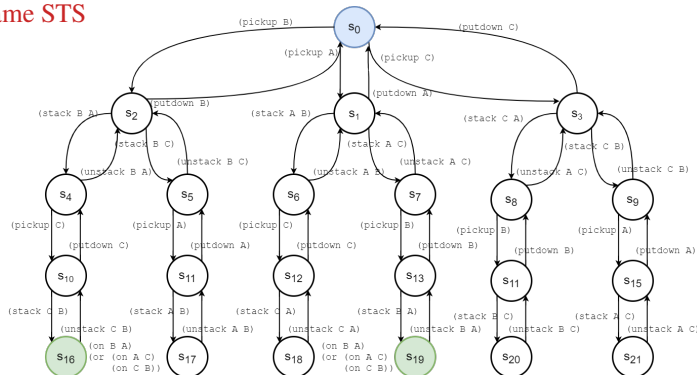
- Adding new actions
 - May lead to previously unreachable states
 - May not result in new solutions at all
 - Still relaxation! Old solutions remain!



RELAXATION EXAMPLE (CONT.)

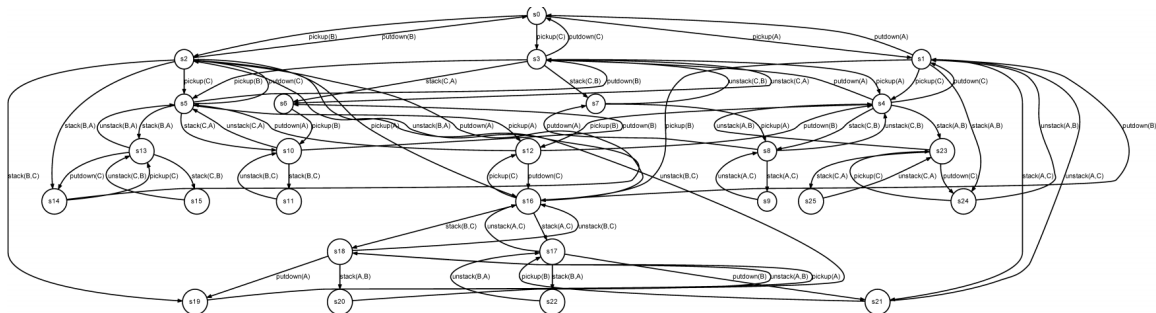
Adding goal states

- New goal formula: $(\text{and } (\text{on } B \ A) \ (\text{or } (\text{on } A \ C) \ (\text{on } C \ B)))$
- All old solutions still valid, but new solutions may exist
- This particular example: Optimal solution from s_0 retains the same length
- Retain the **same STS**




RELAXATION EXAMPLE (CONT.)

- Ignoring state variables
 - Ignore (handempty) fact in preconditions and effects
 - Different state space, no simple addition or removal, but all the old solutions (action sequences) lead from S'_0 to new goal states in S'_g !
 - 22 reachable states \Rightarrow 26
 - 42 transitions \Rightarrow 72



RELAXATION EXAMPLE (CONT.)

● Weakening preconditions of existing actions



- Precondition relaxation: Tiles can be moved across each other
 - Now we have 21 possible first moves: New transitions added to the STS
- All old solutions are still valid, but new ones are added
 - To move "8" into place:
 - Two steps to the right, two steps down, ends up in the same place as "1"

Essentially the same as adding actions!

Result in new transitions!

Can still be solved through search!
 The optimal solution for the relaxed 8-puzzle
 can never be more expensive than the optimal solution for the original 8-puzzle!

RELAXATION HEURISTICS: SUMMARY

- Relaxation: One general principle for designing **admissible** heuristics for **optimal** planning
 - Find a way of transforming planning problems, so that given a problem instance P :
 - Computing its transformation P' is easy (polynomial)
 - Finding an optimal solution to P' is easier than for P
 - All solutions to P are solutions to P' , but the new problem can have additional solutions as well
 - Then the cost of an optimal solution to P' is an admissible heuristic for the original problem P

This is only *one* principle!
There are others, *not* based on relaxation!

SEARCH OR DIRECT COMPUTATION

- As stated:
 - Compute an actual solution π' for the relaxed problem P'
 - Compute $cost(\pi')$
- Example: The 8-puzzle...
 - Ignore (blank ?x ?y) in preconditions and effects
 - Run the problem through an optimal planner
 - Compute the cost of the resulting plan π'

```
(:action move-up
 :parameters (?t ?px ?py ?by)
 :precondition (and
   (tile ?t) (position ?px) (position ?py) (position ?by)
   (dec ?by ?py) (blank ?px ?by) (at ?t ?px ?py))
 :effect (and (not (blank ?px ?by)) (not (at ?t ?px ?py))
   (blank ?px ?py) (at ?t ?px ?by)))
```

SEARCH OR DIRECT COMPUTATION (CONT.)

- But we only use π' to compute its cost!

- Let's analyze the problem (8-tiles) ...

- Each piece has to be moved to the intended row
- Each piece has to be moved to the intended column
- These are **exactly** the required actions given the relaxation!

- \implies **optimal cost** for relaxed problem = sum of Manhattan distances
- \implies **admissible heuristic** for original problem = sum of Manhattan distances
- \implies **Cost** of any optimal solution π' can be computed efficiently *without* π' :

$$\sum_{p \in \text{pieces}} x\text{distance}(p) + y\text{distance}(p)$$

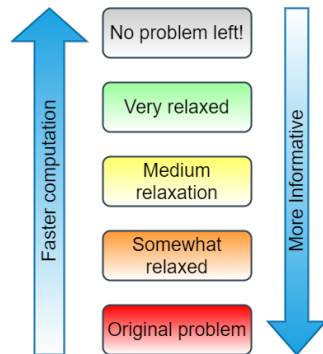
But now we had to **analyze** the problem:

- 1 Decide to ignore "blank"
- 2 Find "sum of manhattan distances"

How do we *automatically* find good relaxations + computation methods? – We will discuss soon!!

RELAXATION HEURISTICS: BALANCE

- The **reason** for relaxation is **rapid calculation**
 - **Shorter solutions** are an *unfortunate side effect*:
Leads to less informative heuristics
 - Relax too much \implies not informative
 - Example: Any piece can teleport into the desired position
 $\implies h(n)$ = number of pieces left to move

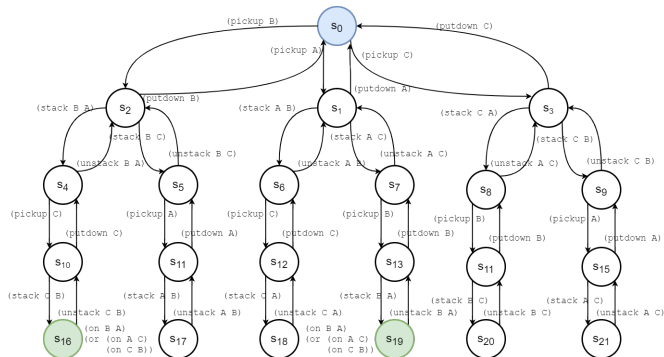


RELAXATION HEURISTICS: IMPORTANT ISSUES!

You **cannot** "use a relaxed problem as a heuristic".

What would that mean?

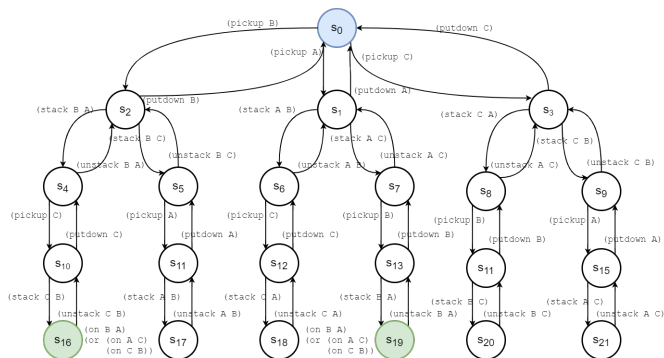
You use the **cost** of an **optimal solution** to the relaxed problem as a heuristic.



This is the problem!

The problem is not a heuristic!

You have to solve it optimally to get the admissibility guarantee.

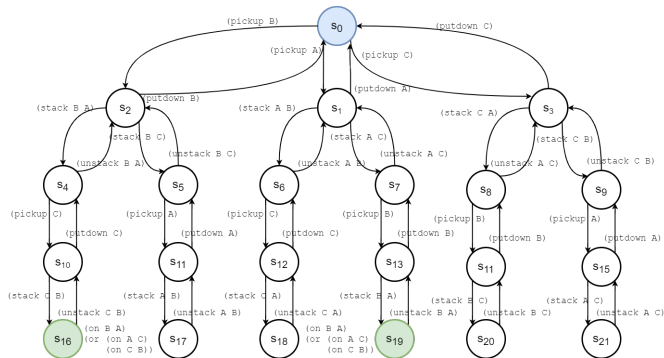


```
(pickup C)
(putdown C)
(pickup B)
(stack B A)
(pickup C)
(stack C B)
```

RELAXATION HEURISTICS: IMPORTANT ISSUES! (CONT.)

You don't just solve the relaxed problem once.

Every time you reach a new state and want to calculate a heuristic, you have to solve the relaxed problem of getting from **that** state to the goal.



Calculate:

$$h(s_0)$$

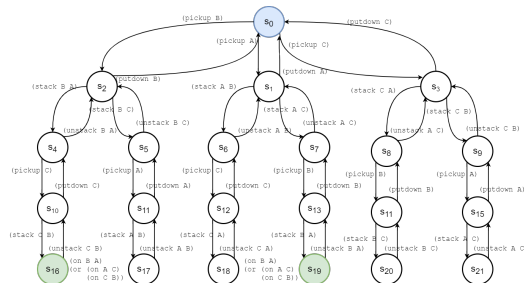
$$h(s_1), h(s_2), h(s_3)$$

... then for every node
you create, depending on
the strategy

-
- The graph illustrates the state space of a 3-disk Tower of Hanoi problem. The root state is s_0 . Transitions are labeled with actions: (pickup A), (putdown A), (stack A B), and (unstack A B). The goal states are s_{16} and s_{19} , which are highlighted in green. The graph shows a complex branching structure with many cycles and dead ends.

RELAXATION HEURISTICS: IMPORTANT ISSUES! (CONT.)

- Relaxation is useful for finding **admissible heuristics**.
- A heuristics cannot be **admissible for some states**.
 - Admissible = does not overestimates cost for **any** state!



REFERENCES I

- [1] Hector Geffner and Blai Bonet. *A Concise Introduction to Models and Methods for Automated Planning*. Synthesis Lectures on Artificial Intelligence and Machine Learning. Morgan & Claypool Publishers, 2013. ISBN 9781608459698. doi: 10.2200/S00513ED1V01Y201306AIM022. URL <https://doi.org/10.2200/S00513ED1V01Y201306AIM022>.
- [2] Malik Ghallab, Dana S. Nau, and Paolo Traverso. *Automated planning - theory and practice*. Elsevier, 2004. ISBN 978-1-55860-856-6.
- [3] Malik Ghallab, Dana S. Nau, and Paolo Traverso. *Automated Planning and Acting*. Cambridge University Press, 2016. ISBN 978-1-107-03727-4. URL <http://www.cambridge.org/de/academic/subjects/computer-science/artificial-intelligence-and-natural-language-processing/automated-planning-and-acting?format=HB>.